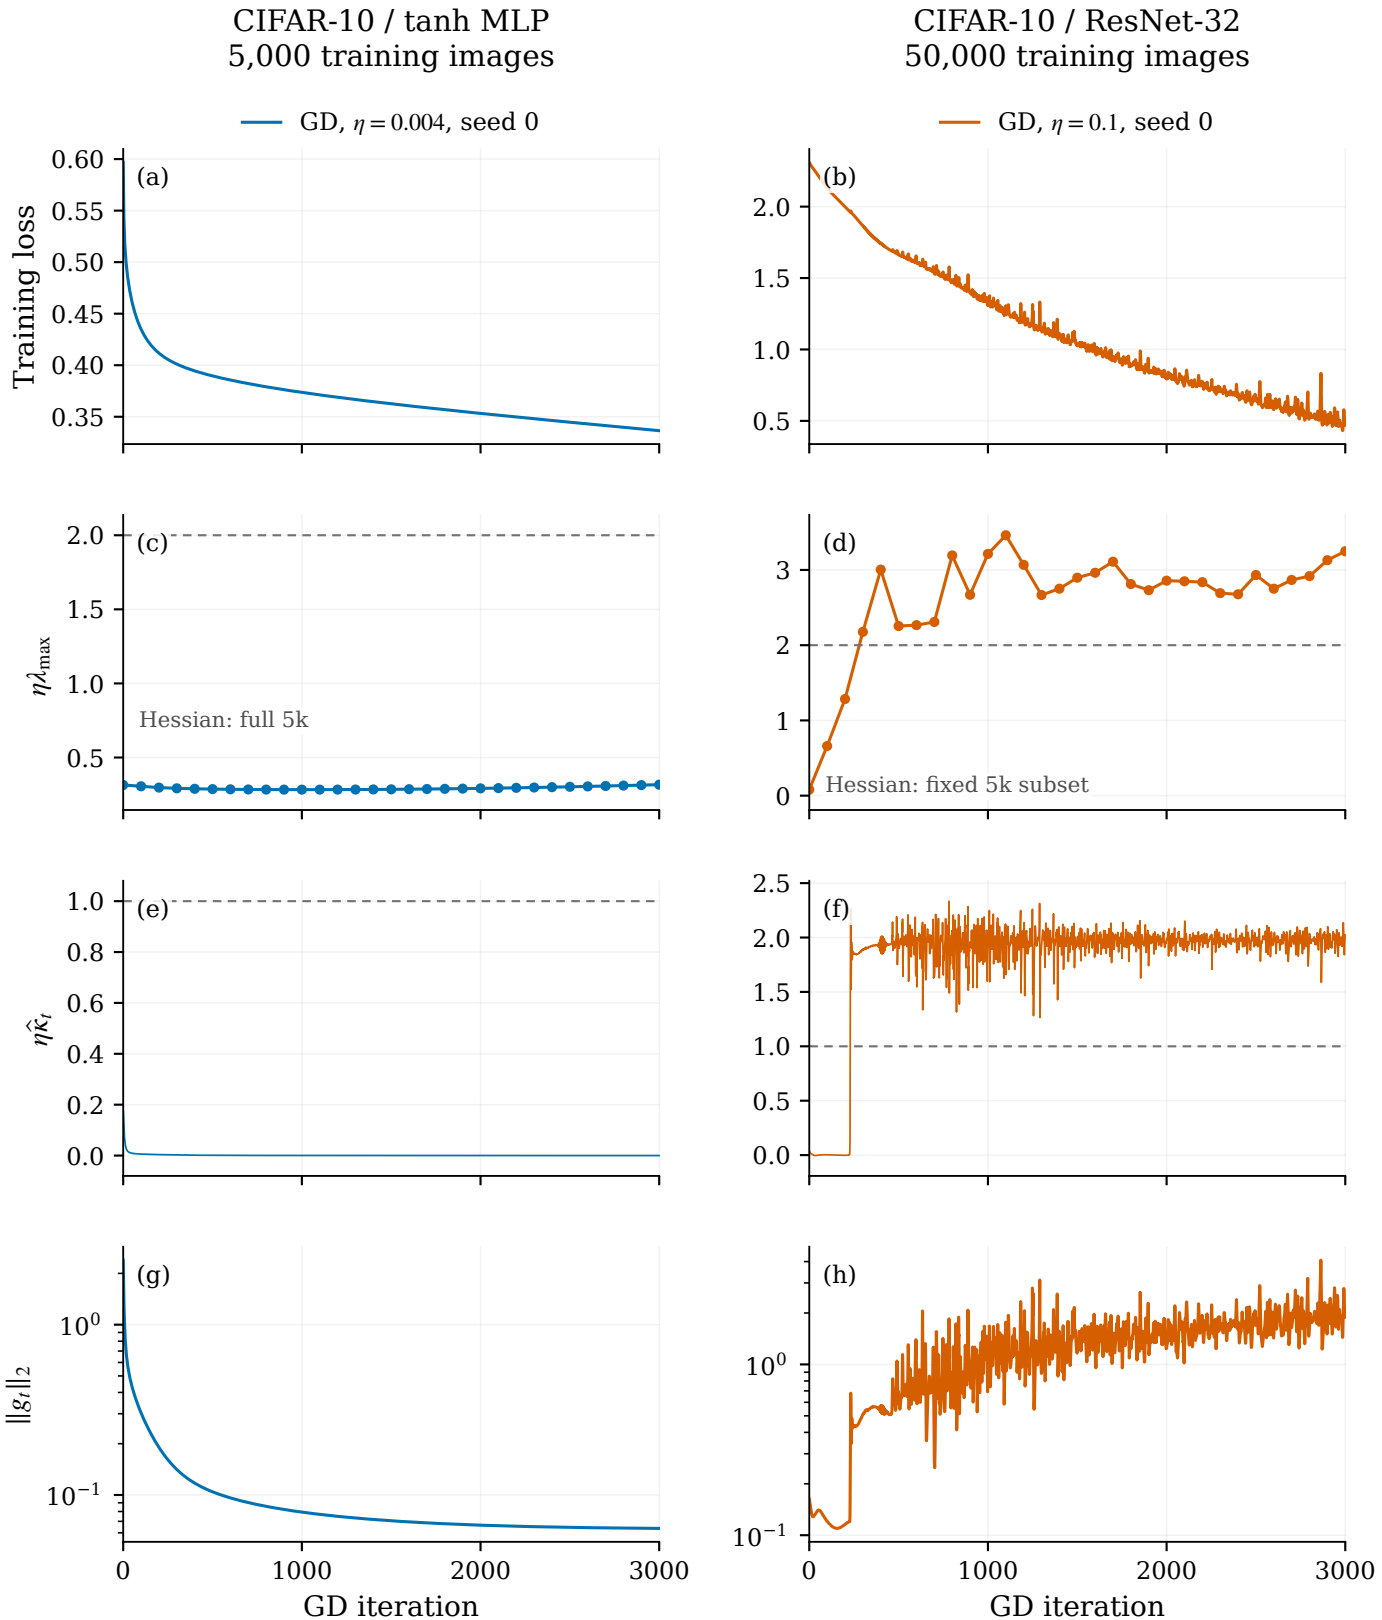

CIFAR-10 / tanh MLP
5,000 training images
CIFAR-10 / ResNet-32
50,000 training images
GD, $\eta = 0.004$, seed 0
GD, $\eta = 0.1$, seed 0
Training loss
(a)
(b)
(c)
(d)
(e)
(f)
(g)
(h)
$\eta\lambda_{\max}$
$\eta\hat{\kappa}_t$
$\|g_t\|_2$
Hessian: full 5k
Hessian: fixed 5k subset
GD iteration
GD iteration