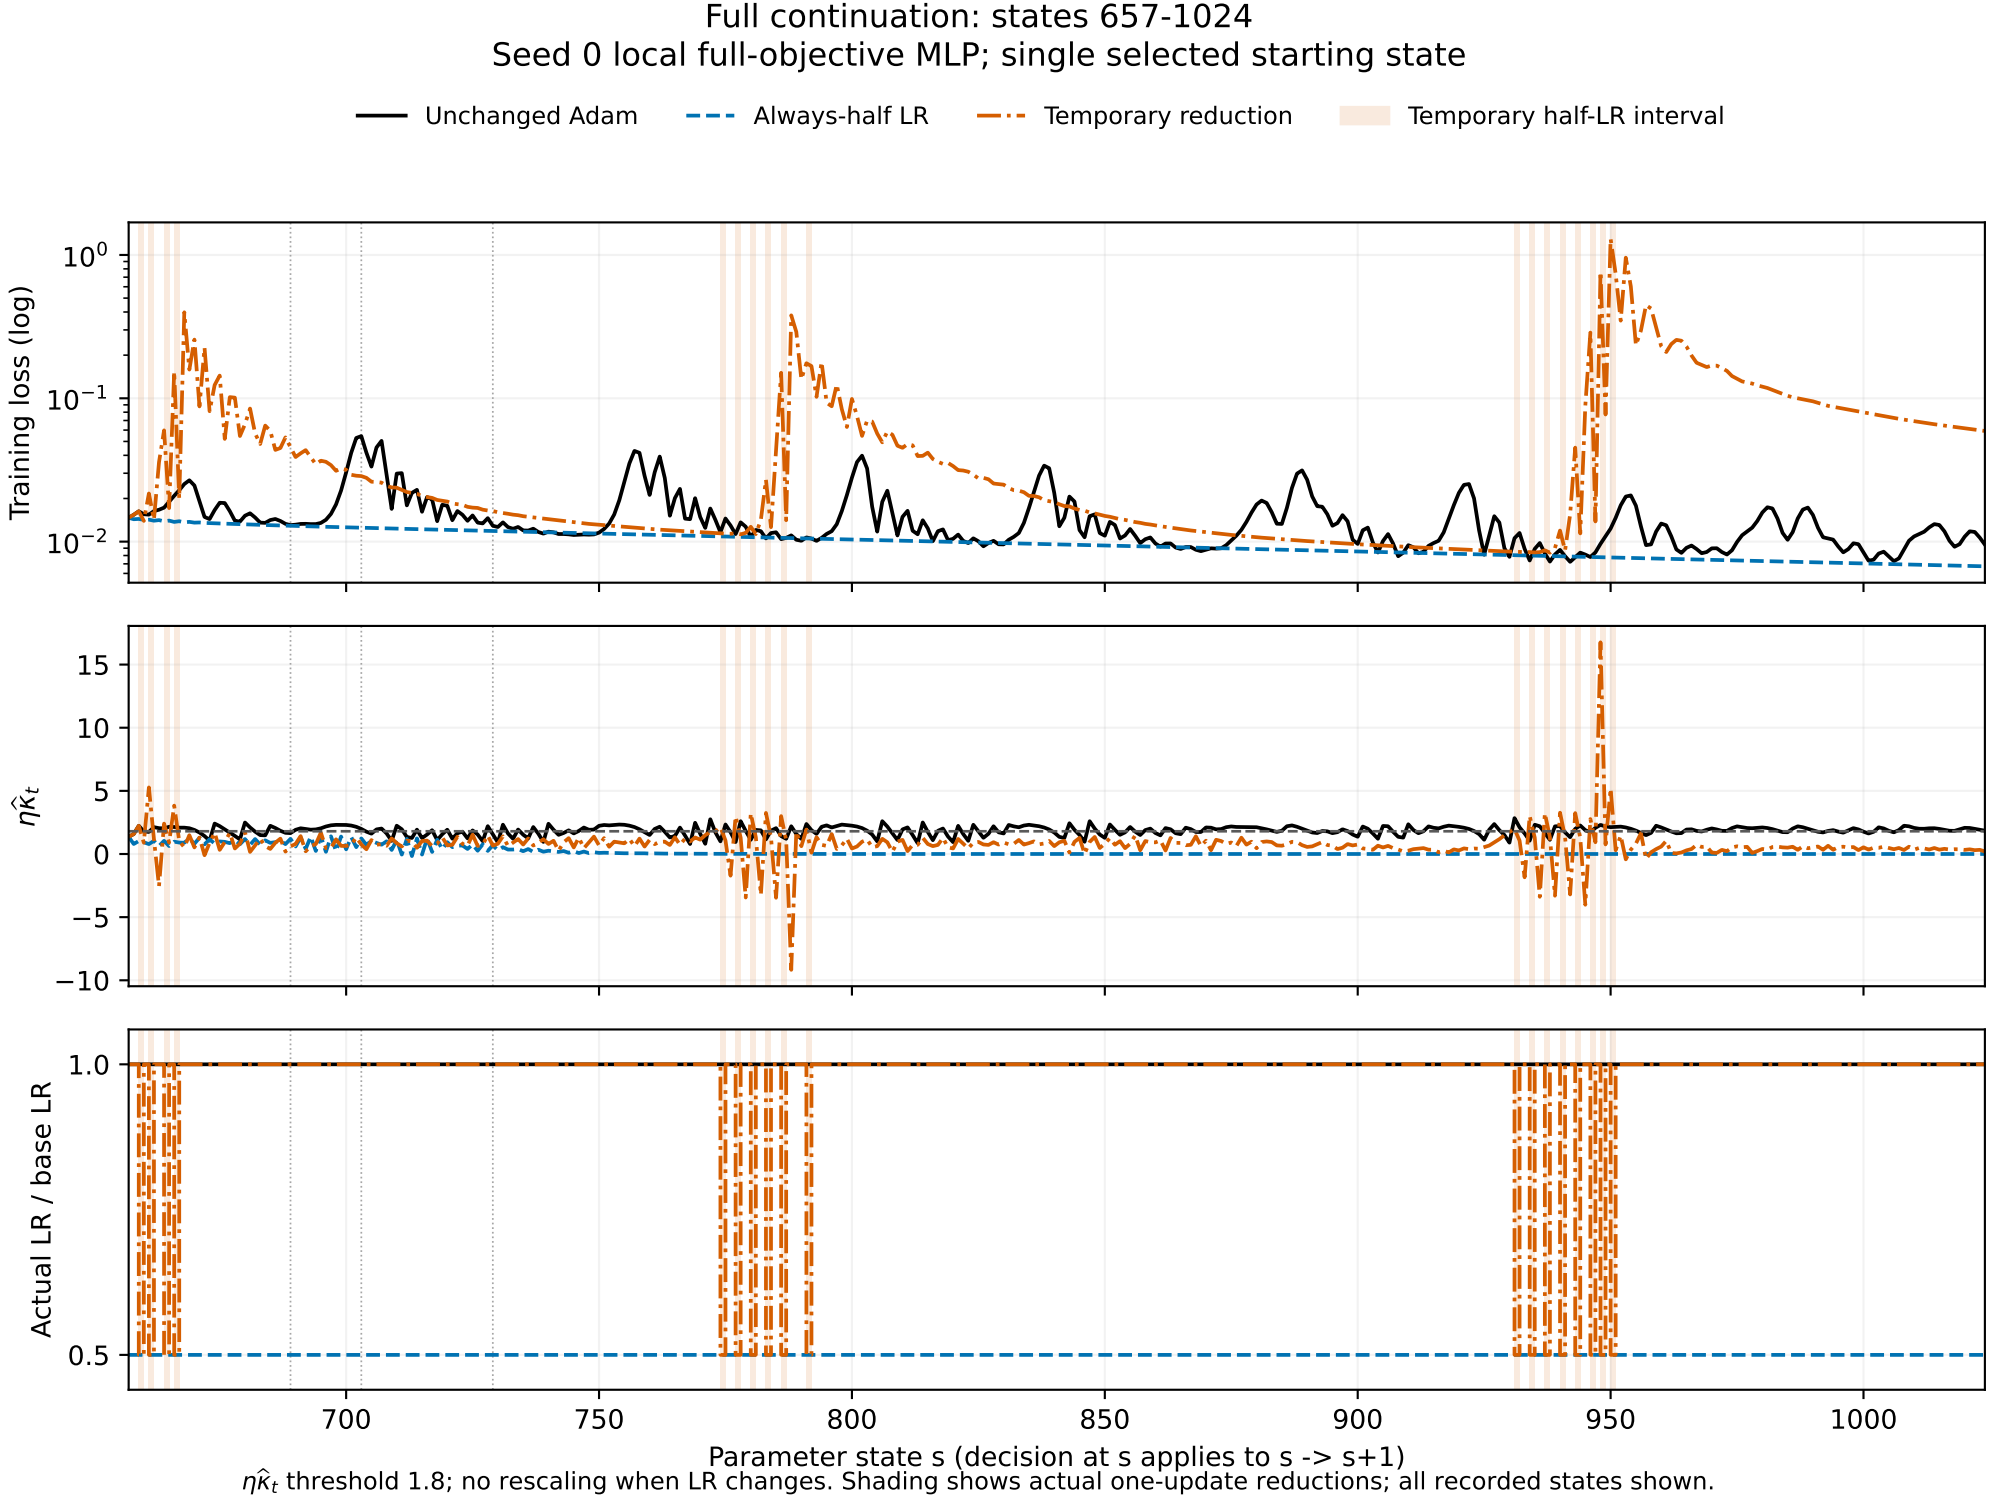
Full continuation: states 657-1024
Seed 0 local full-objective MLP; single selected starting state
Unchanged Adam
Always-half LR
Temporary reduction
Temporary half-LR interval
Training loss (log)
10^{0}
10^{-1}
10^{-2}
$\hat{\eta\kappa}_t$
15
10
5
0
−5
−10
Actual LR / base LR
1.0
0.5
700
750
800
850
900
950
1000
Parameter state s (decision at s applies to s -> s+1)
$\eta\hat{\kappa}_t$ threshold 1.8; no rescaling when LR changes. Shading shows actual one-update reductions; all recorded states shown.